

Tara Dixit

tdixit@stanford.edu | 571-599-1415 | [linkedin.com/in/taradixit](https://www.linkedin.com/in/taradixit) | github.com/taradixit | taradixit.com

EDUCATION

Stanford University — B.S. in Computer Science, AI Track

Graduation Date – Jun 2027

EXPERIENCE

AI Safety & Security Research Intern

Stealth Startup | Jul 2025 – Present

- Ran 60 evaluations across 10 Linux tasks to test how safely AI agents handle setup, debugging, and blocked actions.
- Scored agents on safe commands, verification, unnecessary changes, and recovery from failures.
- Studied guardrail DoS and cyber-agent failures, and designed logs to capture agent behavior.

CURIS Fellow

Stanford University | Jun 2025 – Mar 2026

- Analyzed 31 threat reports from incident-response firms and cyber insurers to compare how they describe real-world attacks and prioritize threats.
- Found major disagreement on initial access and exploited vulnerabilities, limiting their usefulness for security teams.

Research Assistant — Empirical Security Research Group

Stanford University | Jan 2025 – Jun 2025

- Contributed to a USENIX Security 2025 study of cybersecurity guidance across 41 countries.
- Built a comparison taxonomy and used LLMs to compare recommendations, finding inconsistencies that could create false confidence about security coverage.

PROJECTS

ContainmentBench — Security Boundaries for Coding Agents

Jun 2026 – Present

- Built a Linux containment testbed that limits coding agents' file, process, and network access during multi-step tasks.
- Instrumented the environment to detect and trace attempted policy violations.

Hidden Reliability Failures in Malware Detection

Mar 2026 – Aug 2026

- Evaluated the EMBER2024 malware detector on 540,000 Windows executable files.
- Despite 98.0% overall accuracy and strong overall calibration, it missed 85.2% of samples from one malware family, showing that strong aggregate results hid severe family-level failures.

Predicting Vulnerability Attributes from CVE Descriptions

Mar 2026 – Aug 2026

- Built models to predict vulnerability attributes from CVE security advisories, trained on 2022–2023 CVEs and evaluated on 36,539 CVEs from 2024.
- Reached 86.1% average accuracy across CVSS components and found that strong overall performance did not extend to rare and previously unseen labels.

PUBLICATIONS

- Investigating Usefulness of Security Threat Reports — manuscript under review
- Privacy and Paternalism: The Ethics of Student Data Collection — MIT Schwarzman College of Computing (2022)

LEADERSHIP

President, Women in Applied Cyber Club

Stanford University | 2025 – Present

- Led cybersecurity workshops and organized talks with industry professionals.

TECHNICAL SKILLS

Programming: Python, C++

ML: PyTorch, scikit-learn, LightGBM, Hugging Face Transformers, pandas, NumPy

Systems/Tools: Linux, Git, Bash, Docker